

# HYBRID FORTRAN

High Performance, Low Friction GPGPU for Weather Prediction

Michel Müller

RIKEN Advanced Institute for Computational Science  
michel@typhooncomputing.com

Supervised by

Dr. Takashi Shimokawabe  
Tokyo Institute of Technology

Prof. Dr. Takayuki Aoki  
Tokyo Institute of Technology

Prof. Dr. Roger Wattenhofer  
ETH Zurich

Dr. Tabito Hara  
Japan Meteorological Agency

Dr. Naoya Maruyama  
RIKEN Advanced Institute for Computational Science



東京工業大學  
Tokyo Institute of Technology

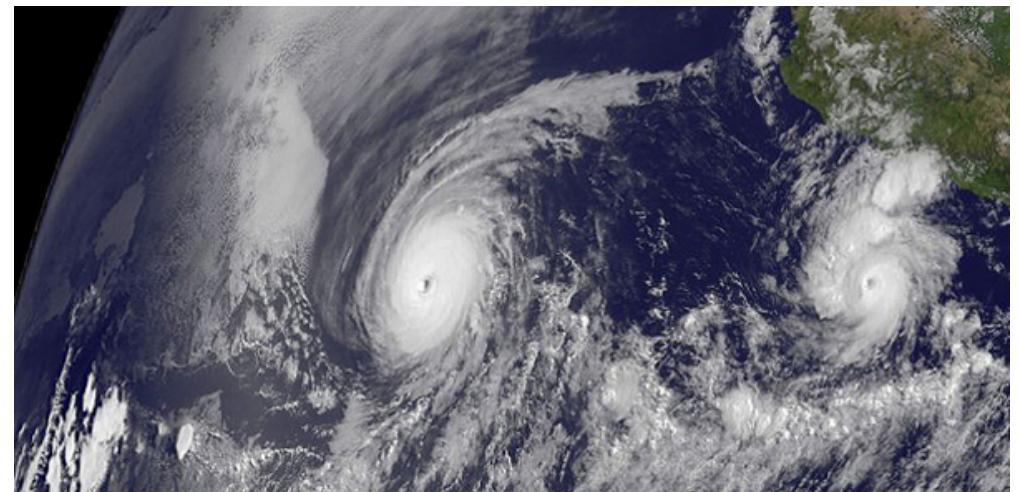


ETH

Eidgenössische Technische Hochschule Zürich  
Swiss Federal Institute of Technology Zurich



Japan Meteorological  
Agency



Creative Commons: Nasa Goddard Space Flight Centre, 2010

# HYBRID FORTRAN

High Performance, Low Friction GPGPU for Weather Prediction

CPU + GPU executable

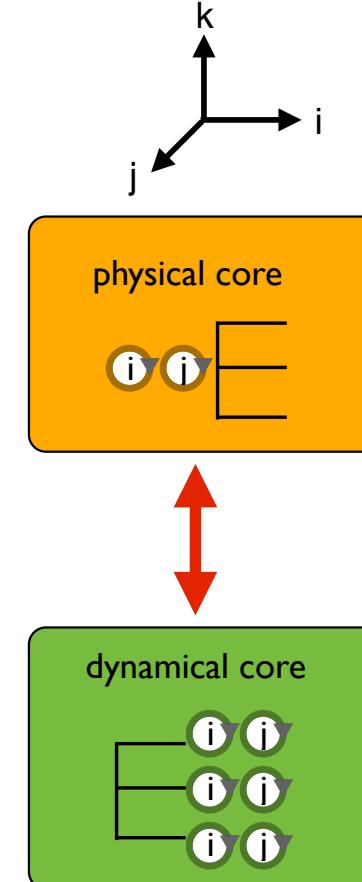
## Topics |

1. Motivation
2. Why not OpenACC?
3. Introducing Hybrid Fortran
4. Implementation
5. Results & Discussion
6. Conclusion & Future Work

# ASUCA

Motivation

- What is ASUCA? [4]
  - Nonhydrostatic weather prediction model
  - Successor of current JMANHM model
  - Dynamical + physical core
- Goals for ASUCA GPGPU portation of physical core:
  1. Eliminate host-to-device communication
  2. Gain execution time speedups



[1] K. Kawano, J. Ishida, and C. Muroi, “Development of a New Nonhydrostatic Model ASUCA at JMA”, 2010.



# ASUCA: Challenges

Motivation

- Challenges:
  1. Keep familiar development environment  
(Fortran 90)
  2. Amount of code changes as small as possible
  3. Keep CPU compatibility
  4. Keep CPU performance and enable high GPU performance
- Scope of this Project:
  - Concentrate on single GPU / single CPU socket

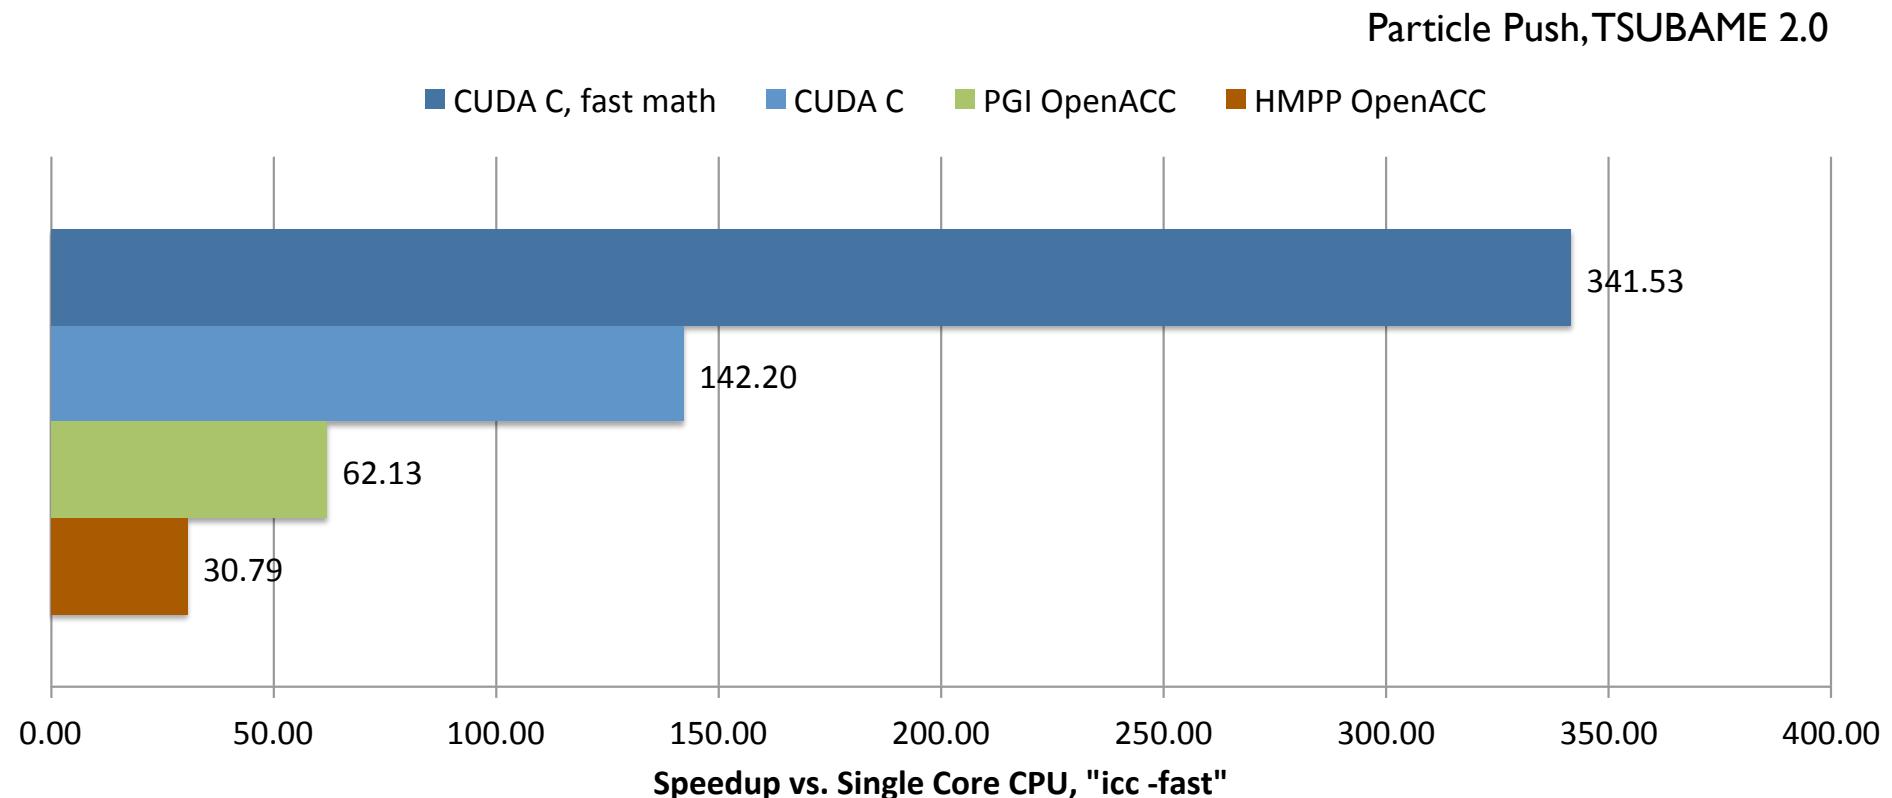


Why not OpenACC?



## Why not OpenACC?

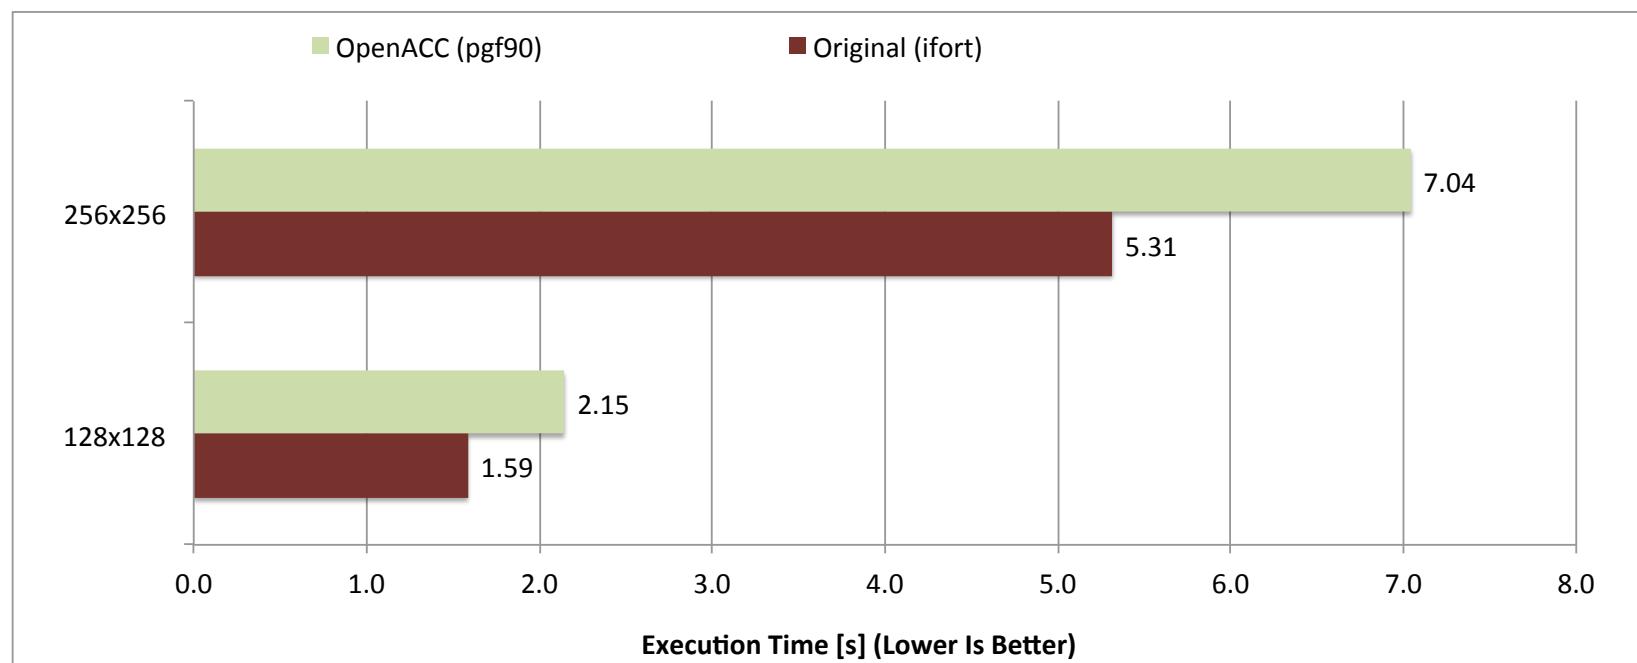
### 1. GPU performance for computationally bounded problems



## Why not OpenACC?

1. GPU performance for computationally bounded problems
2. CPU performance

Shortwave Radiation (8 Kernels), CPU 6 Core, TSUBAME 2.0

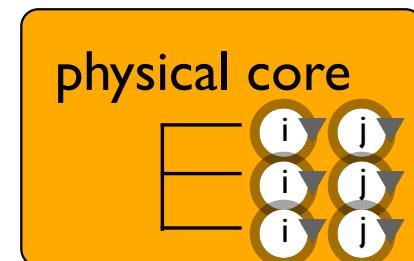
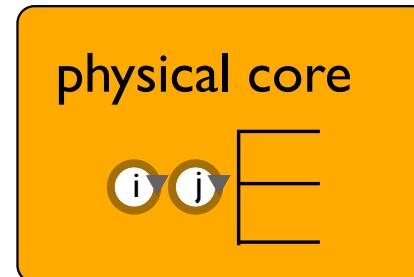


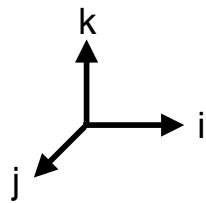
## Why not OpenACC?

1. GPU performance for computationally bounded problems
2. CPU performance
3. Insufficient abstraction

Original  
(CPU)

GPU  
optimized



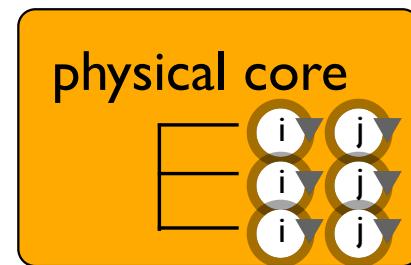
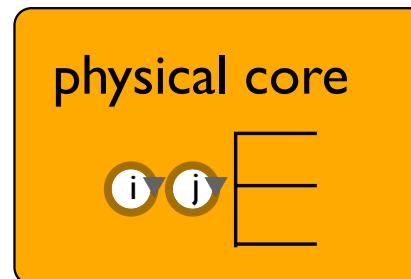


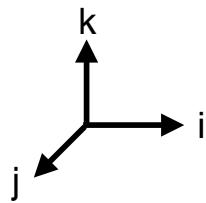
Hybrid Fortran  
approach

Original  
(CPU)

GPU  
optimized

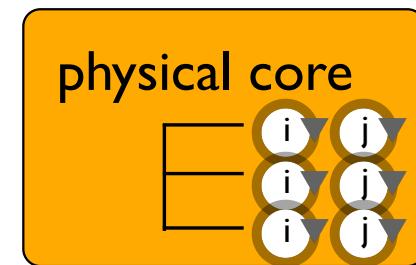
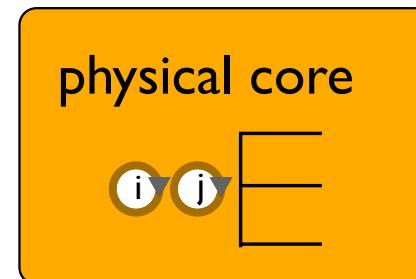
## Introduction | Hybrid Fortran





Hybrid Fortran  
approach

## Introduction | Hybrid Fortran



# Example | Hybrid Fortran

```
1 module example
2 contains
3     subroutine wrapper(a, b, c, d)
4         real, intent(in) :: a(DOM(NX, NY, NZ)), b(DOM(NX, NY, NZ))
5         real, intent(out) :: c(DOM(NX, NY, NZ)), d(DOM(NX, NY, NZ))
6         integer(4) :: x, y
7
8         do y=1,NY
9             do x=1,NX
10                call add(a(AT(x,y,:)), b(AT(x,y,:)), c(AT(x,y,:)))
11                call mult(a(AT(x,y,:)), b(AT(x,y,:)), d(AT(x,y,:)))
12            end do
13        end do
14    end subroutine
15
16    subroutine add(a, b, c)
17        real, intent(in) :: a(NZ), b(NZ)
18        real, intent(out) :: c(NZ)
19        integer :: z
20
21        do z=1,NZ
22            c(z) = a(z) + b(z)
23        end do
24    end subroutine
```

Motivation

OpenACC

Hybrid Fortran

Implementation

Discussion

Conclusion

# Example | Hybrid Fortran

```
1 module example
2 contains
3   subroutine wrapper(a, b, c, d)
4     real, intent(in) :: a(DOM(NX, NY, NZ)), b(DOM(NX, NY, NZ))
5     real, intent(out) :: c(DOM(NX, NY, NZ)), d(DOM(NX, NY, NZ))
6     integer(4) :: x, y
7
8     do y=1,NY
9       @parallelRegion{appliesTo(CPU), domName(x,y), domSize(NX, NY)}
10      call add(a(AT(x,y,:)), b(AT(x,y,:)), c(AT(x,y,:)))
11      call mult(a(AT(x,y,:)), b(AT(x,y,:)), d(AT(x,y,:)))
12    @end parallelRegion
13    end do
14  end subroutine
15
16  subroutine add(a, b, c)
17    real, intent(in) :: a(NZ), b(NZ)
18    real, intent(out) :: c(NZ)
19    integer :: z
20
21    do z=1,NZ
22      c(z) = a(z) + b(z)
23    end do
24  end subroutine
```

Motivation

OpenACC

Hybrid Fortran

Implementation

Discussion

Conclusion

# Example | Hybrid Fortran

```
1 module example
2 contains
3   subroutine wrapper(a, b, c, d)
4     real, intent(in) :: a(DOM(NX, NY, NZ)), b(DOM(NX, NY, NZ))
5     real, intent(out) :: c(DOM(NX, NY, NZ)), d(DOM(NX, NY, NZ))
6     integer(4) :: x, y
7
8     do y=1,NY
9       @parallelRegion{appliesTo(CPU), domName(x,y), domSize(NX, NY)}
10      call add(a(AT(x,y,:)), b(AT(x,y,:)), c(AT(x,y,:)))
11      call mult(a(AT(x,y,:)), b(AT(x,y,:)), d(AT(x,y,:)))
12    @end parallelRegion
13    end do
14  end subroutine
15
16  subroutine add(a, b, c)
17    real, intent(in) :: a(NZ), b(NZ)
18    real, intent(out) :: c(NZ)
19    integer :: z
20
21    @parallelRegion{appliesTo(GPU), domName(x,y), domSize(NX, NY)}
22    do z=1,NZ
23      c(z) = a(z) + b(z)
24    end do
25    @end parallelRegion
26  end subroutine
27
28
```

Motivation

OpenACC

Hybrid Fortran

Implementation

Discussion

Conclusion

# Example | Hybrid Fortran

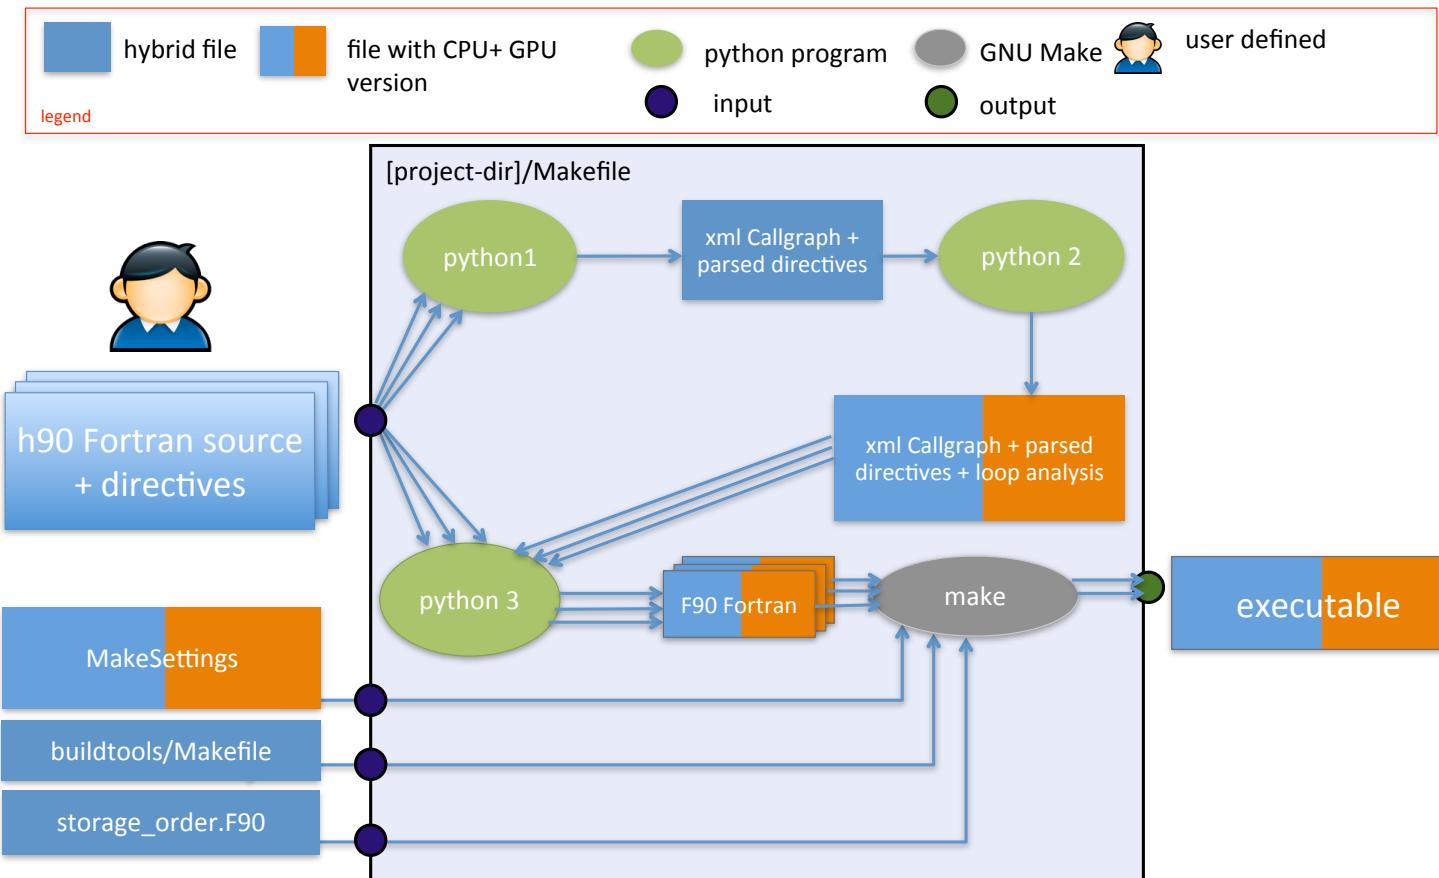
```
1 module example
2 contains
3   subroutine wrapper(a, b, c, d)
4     real, dimension(NZ), intent(in) :: a, b
5     real, dimension(NZ), intent(out) :: c, d
6
7     @domainDependant{domName(x,y,z), domSize(NX,NY,NZ), domPP(DOM),
8       accPP(AT)}
9     a, b, c, d
10    @end domainDependant
11
12    @parallelRegion{appliesTo(CPU), domName(x,y), domSize(NX, NY)}
13    call add(a, b, c)
14    call mult(a, b, d)
15    @end parallelRegion
16  end subroutine
```



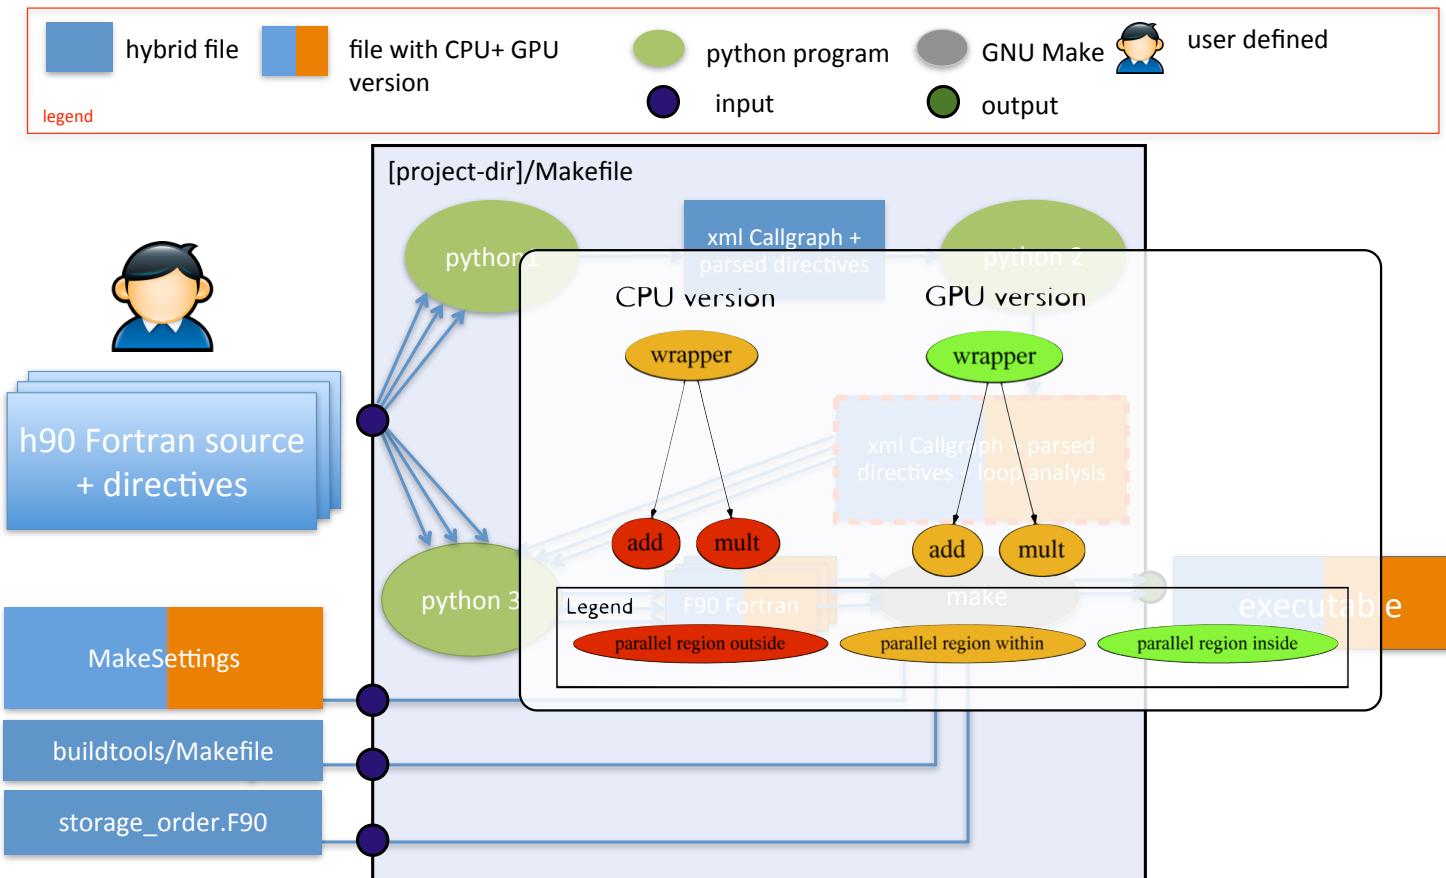
## Implementation



# Build System | Implementation



# Build System | Implementation



Motivation

OpenACC

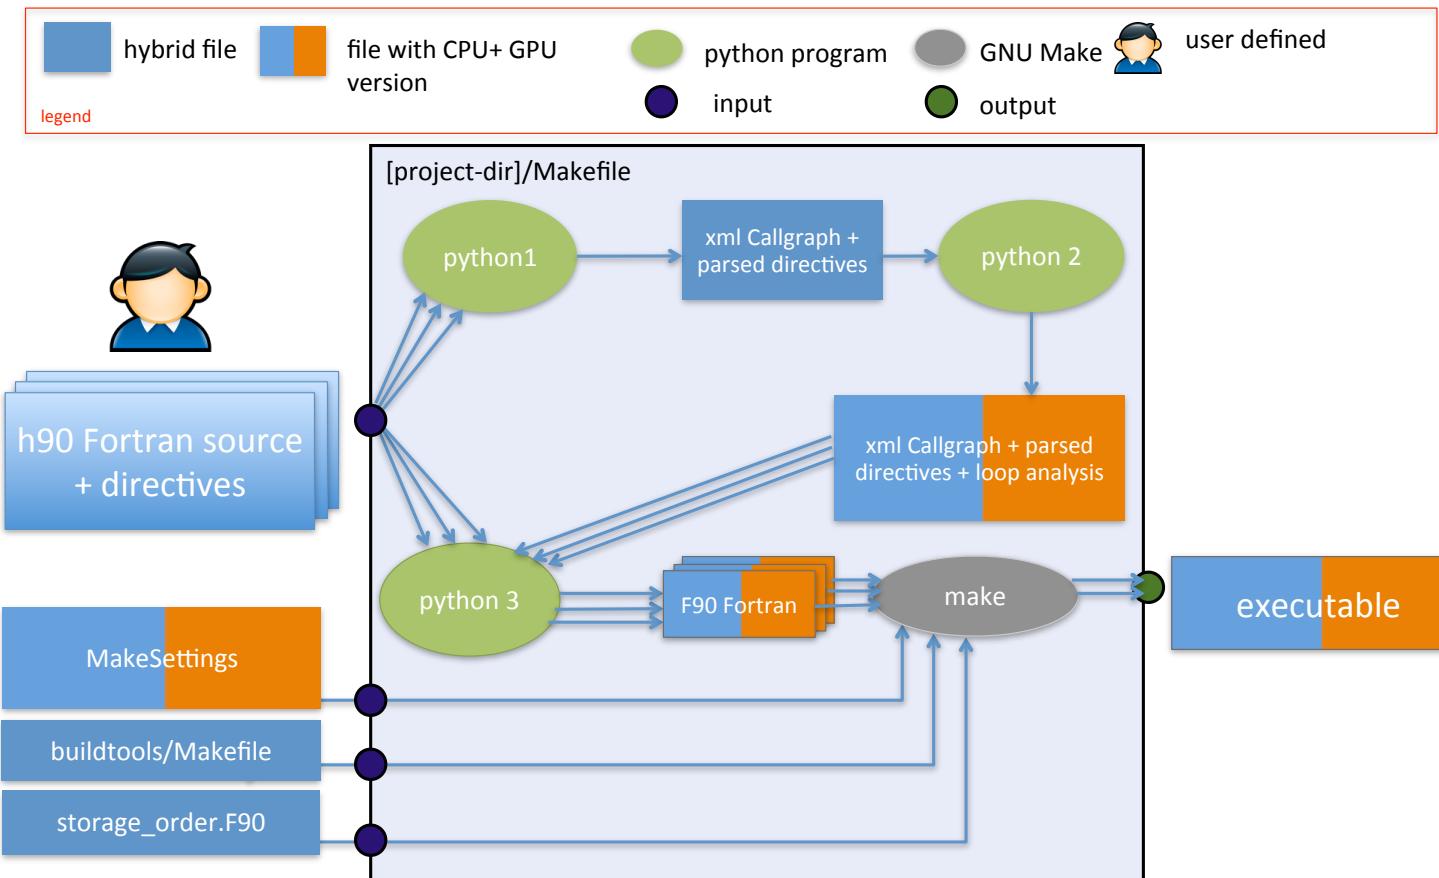
Hybrid Fortran

Implementation

Discussion

Conclusion

# Build System | Implementation



Motivation

OpenACC

Hybrid Fortran

Implementation

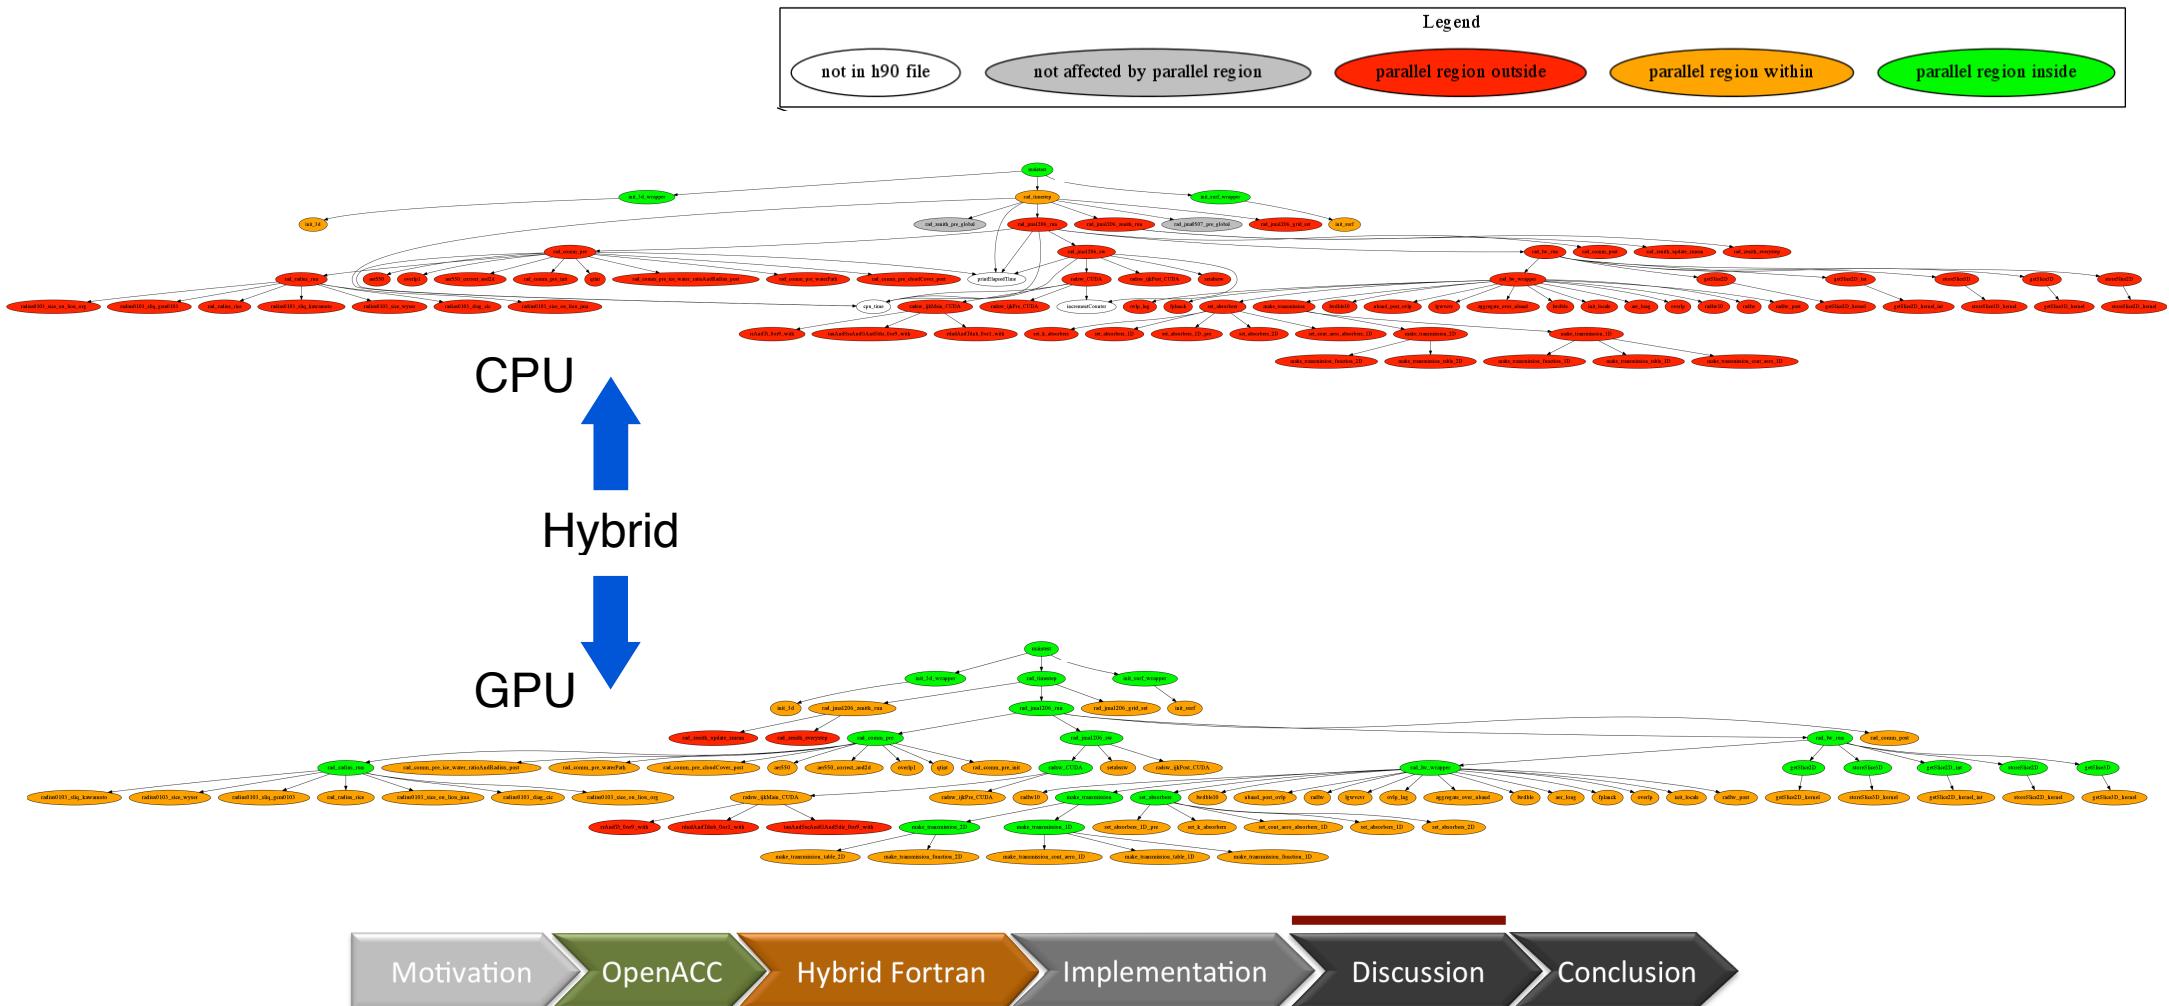
Discussion

Conclusion

## Results & Discussion

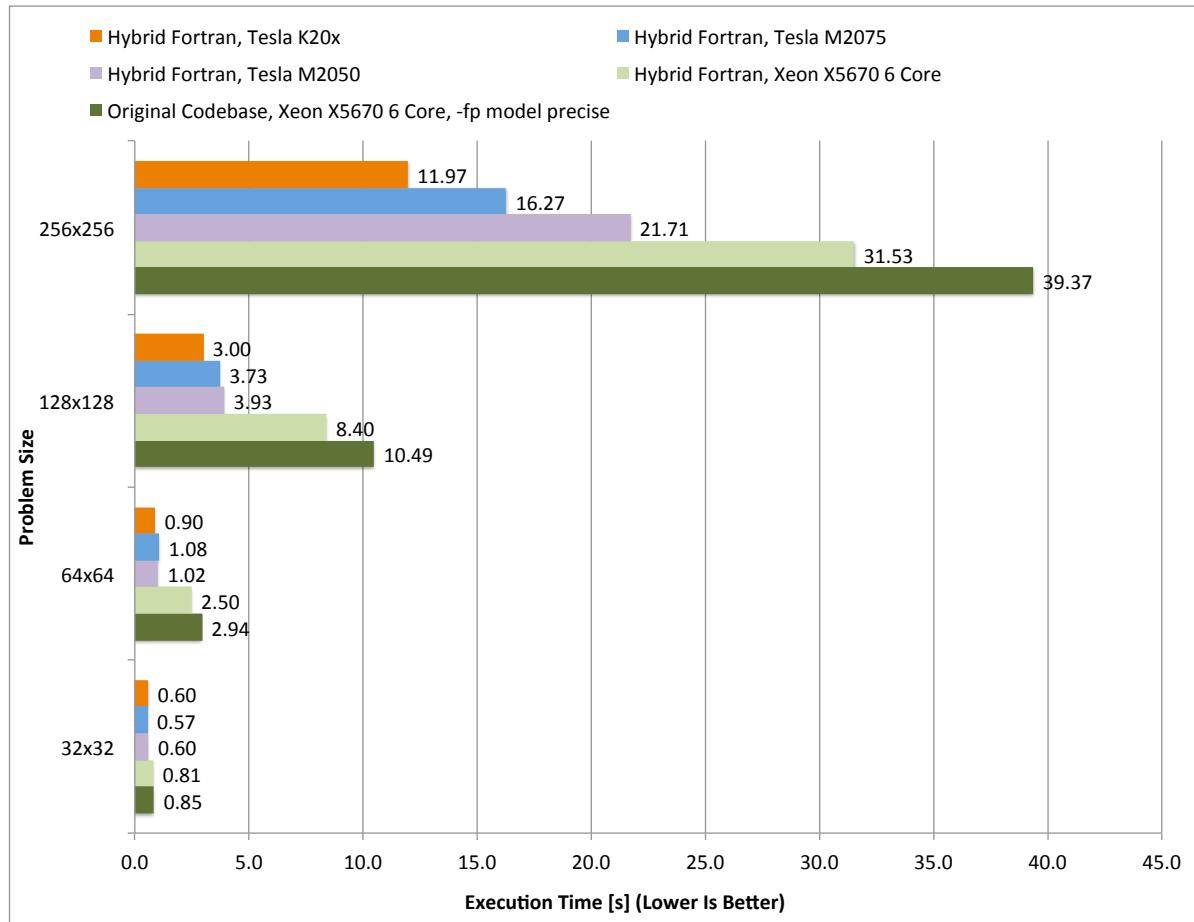


# Scope of Current Implementation | Results & Discussion



# Results: Execution Time

## Results & Discussion



Motivation

OpenACC

Hybrid Fortran

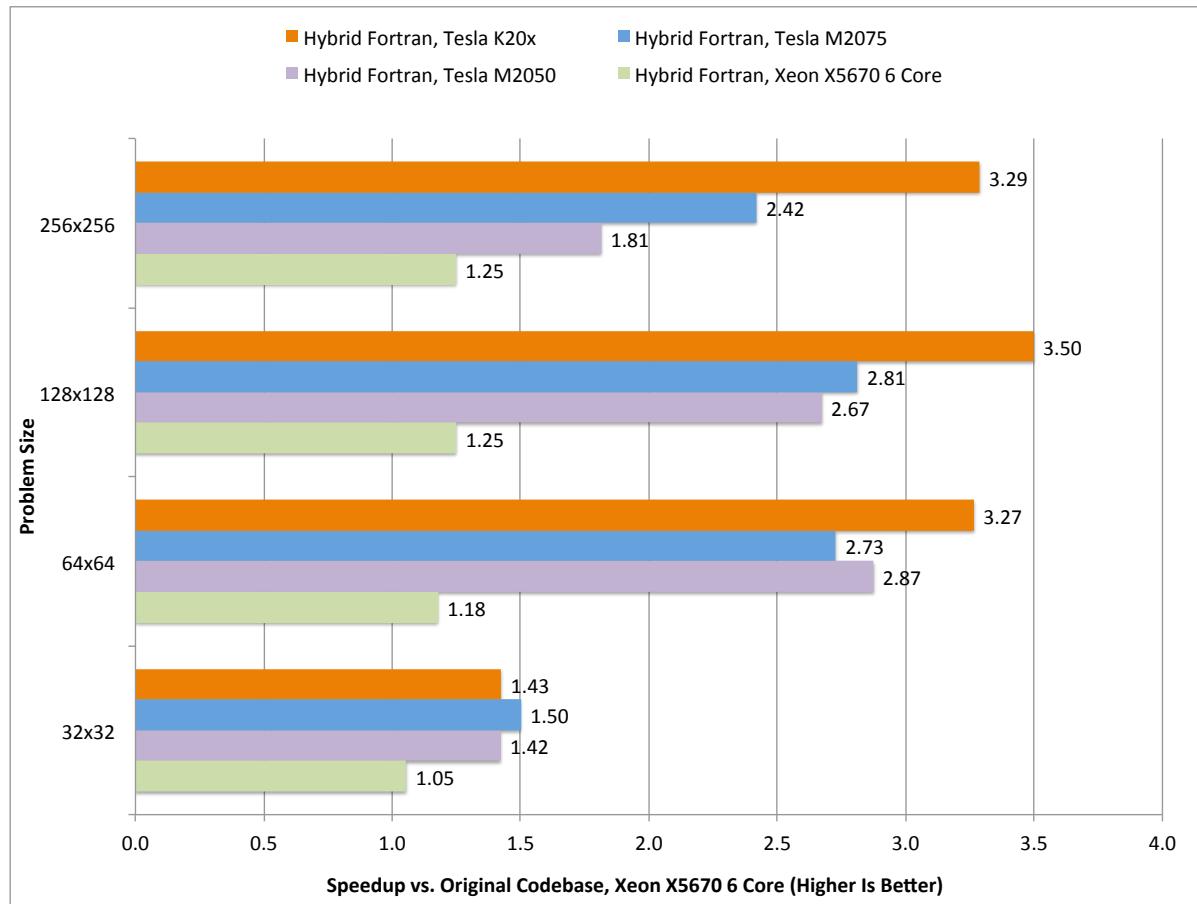
Implementation

Discussion

Conclusion

# Results: Speedup

## Results & Discussion



# Expected Speedup | Results & Discussion



Tesla  
M2050

vs.



Tesla  
M2050

vs.

Xeon  
X5670

Memory Bandwidth Bounded Algorithms *	11.1 x	5.3 x
Computationally Bounded Algorithms [2, p. 4], [3], [5]	44 x	7.4 x

\* Tesla sustained bandwidth according to CUDA SDK bandwidth test; CPU bandwidths according to [2, p. 4], [4, p. 7].

Motivation

OpenACC

Hybrid Fortran

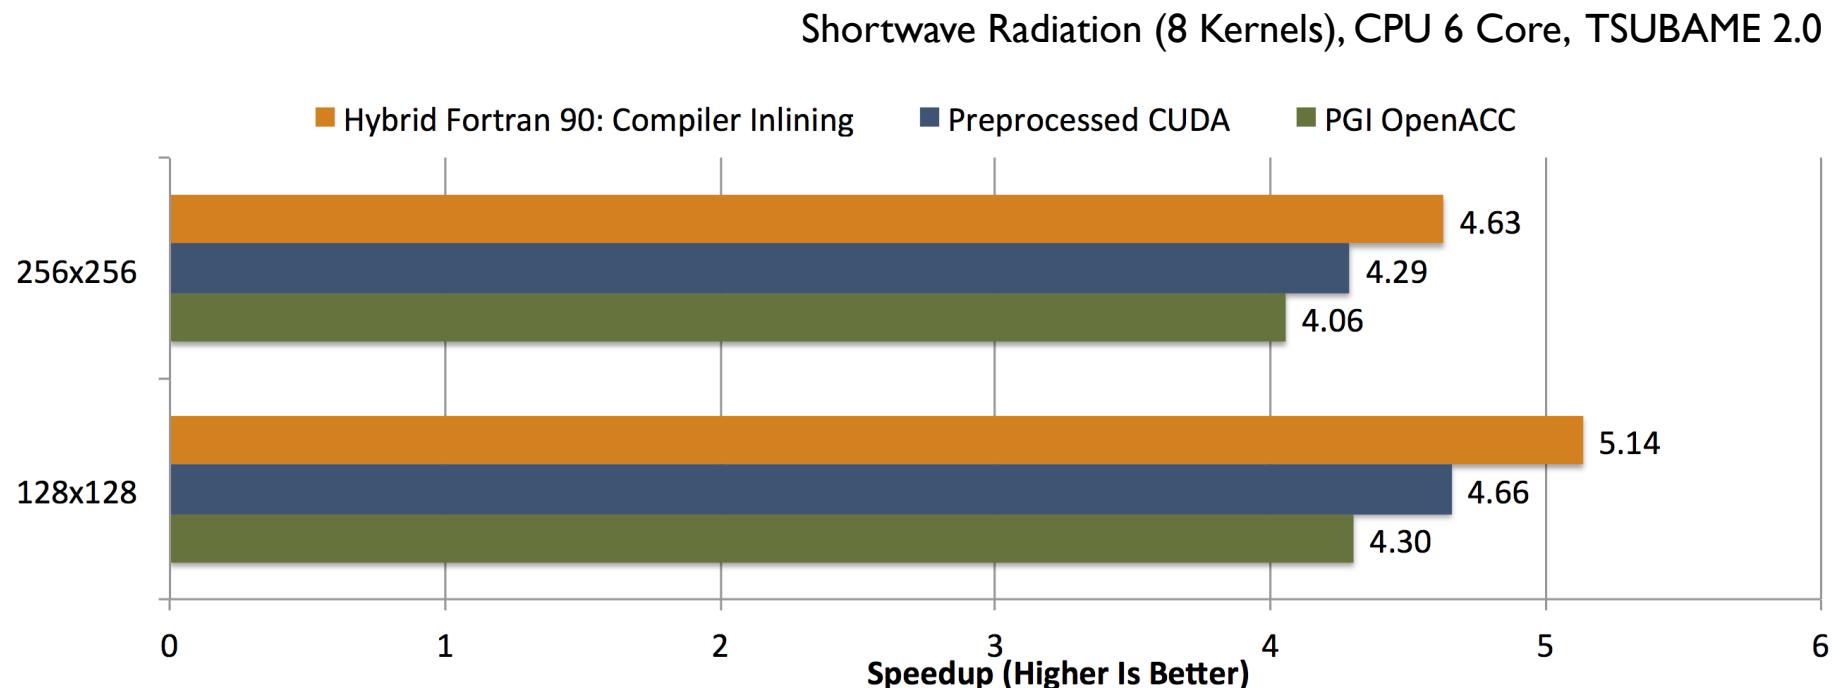
Implementation

Discussion

Conclusion

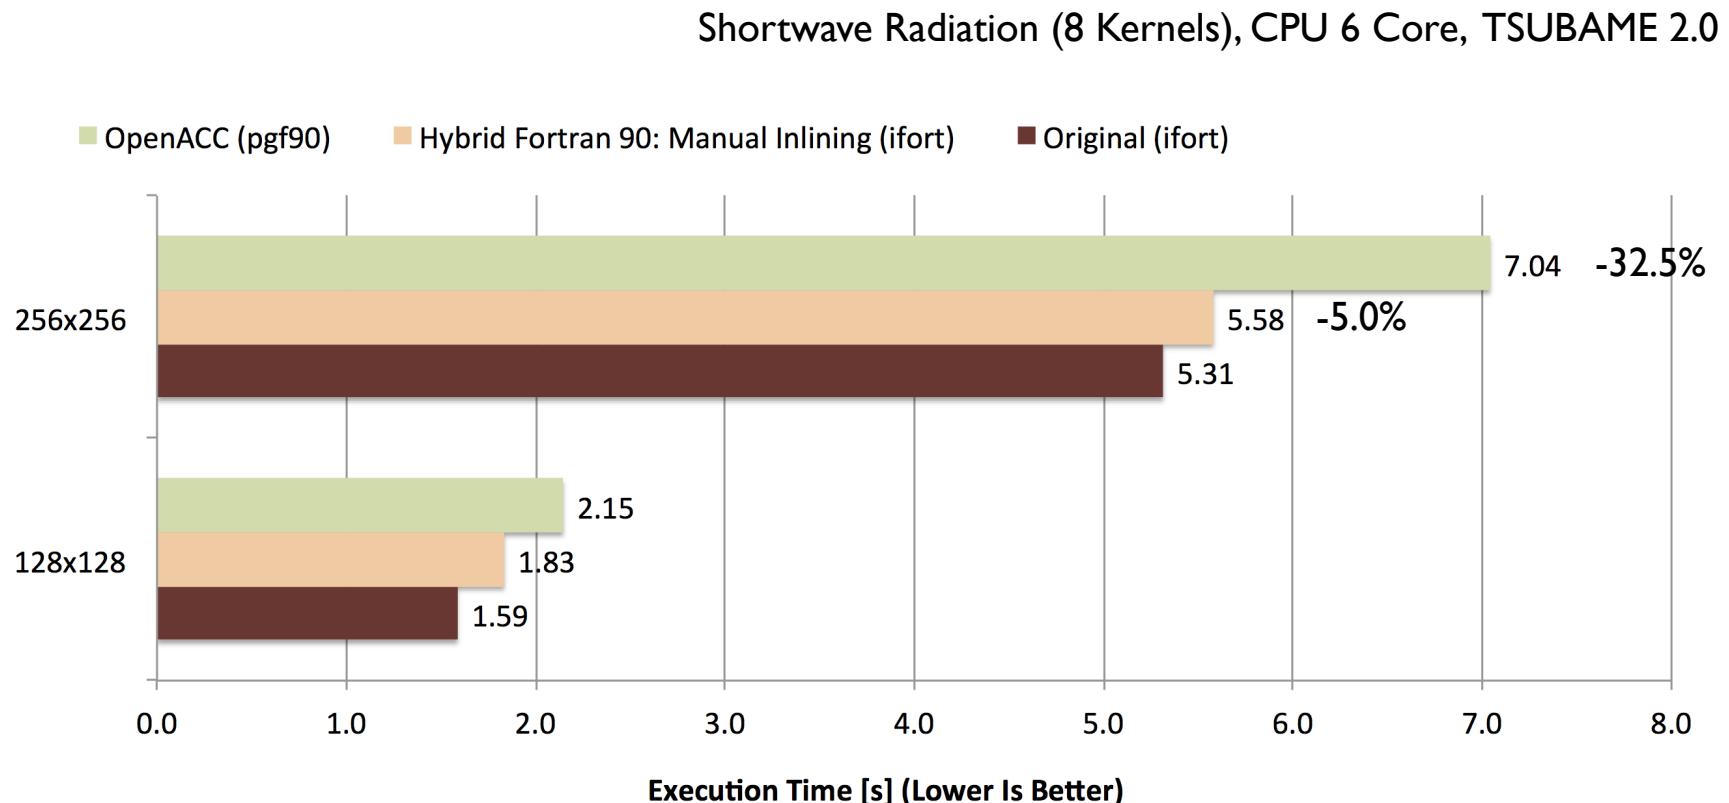
# Tuned Results: GPU

Results & Discussion

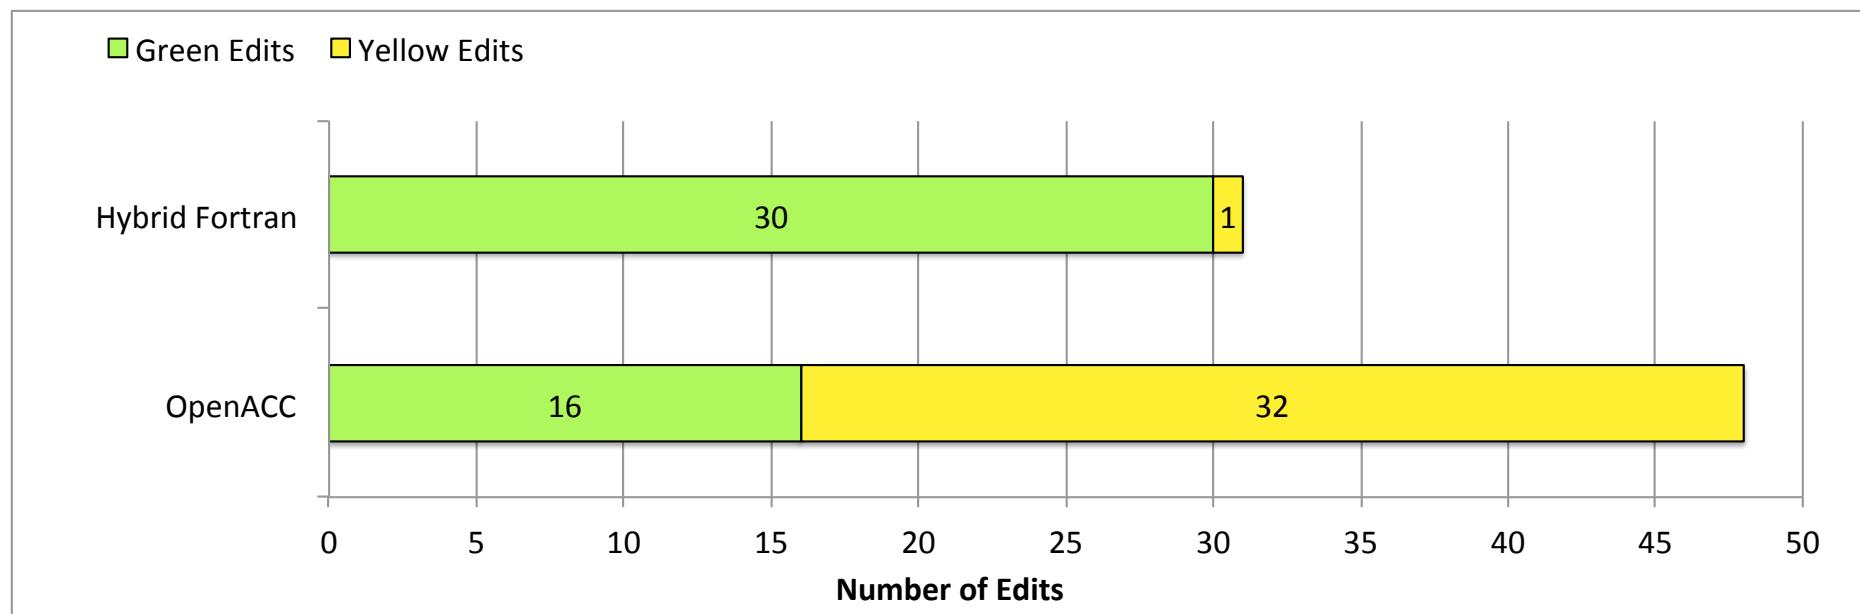


# Tuned Results: CPU

Results & Discussion



# Usability | Results & Discussion



# Conclusion

1. Dynamic scripting languages can be successfully used to enhance the functionality of a compiled language.
2. Hybrid Fortran allows us to get GPU compatibility with baseline performance quickly, without compromising CPU performance.
3. We still have all the flexibility for performance tuning.



## Future Work

- Full scale ASUCA Physical Core implementation
- Pending Improvements for **Hybrid Fortran**
  - 1.General Stencil compatibility
  - 2.Support for more implementations and hardware, e.g. OpenCL, Xeon Phi
  - 3.Data handling simplifications

## References

- [1] K. Kawano, J. Ishida, and C. Muroi, “Development of a New Nonhydrostatic Model ASUCA at JMA”, 2010.
- [2] Intel Corporation, “Specifications Intel Xeon Processor X5670, [http://ark.intel.com/products/47920/Intel-Xeon-Processor-X5670-% 2812M-Cache-2% 93-GHz-6% 2040-GTs-Intel-QPI%29](http://ark.intel.com/products/47920/Intel-Xeon-Processor-X5670-%202812M-Cache-2%2093-GHz-6%2040-GTs-Intel-QPI%29)”, 2010. [Online; accessed October 19, 2012]
- [3] A. Vladimirov, “Arithmetics on Intel’s Sandy Bridge and Westmere CPUs”, [http://research.colfaxinternational.com/file.axd?file=2012%2F4%2FColfax\\_FLOPS.pdf](http://research.colfaxinternational.com/file.axd?file=2012%2F4%2FColfax_FLOPS.pdf) 2012. ; [Online, accessed October 19, 2012].
- [4] NVIDIA, “Specifications NVIDIA Tesla C2050 GPU computing processor”, [http://www.nvidia.com/docs/IO/43395/NV\\_DS\\_Tesla\\_C2050\\_C2070\\_jul10\\_lores.pdf](http://www.nvidia.com/docs/IO/43395/NV_DS_Tesla_C2050_C2070_jul10_lores.pdf), 2010. [Online; accessed September 26, 2012]
- [5] S. Williams, A. Waterman, and D. Patterson, “Roofline: An Insightful Visual Performance Model for Multicore Architectures”, 2009.

## Next Steps + Questions |

- Hybrid Fortran will be available under LGPL license by this weekend!
- E-Mail me for GitHub-link, thesis, slides, further questions..

[michel@typhooncomputing.com](mailto:michel@typhooncomputing.com)